

Name: Anshul Patni
Name: Prasun Saxena

Student ID: 012415156
Student ID: 012449775

Question 1: For each member in your team, provide 1 paragraph detailing what parts of the lab that member implemented / researched.

Answer:

1. Work done by Anshul Patni:

- Downloaded and built the Linux Kernel modules and associated libraries to create a local copy of Linux Kernel.
- Modified the cmpe283-1.c code by adding the custom logic to enable our system to read and give output for capabilities of the various MSRs.
- Staged and committed the cmpe283-1.c code file and Makefile after inserting the module and printing out the buffer from the Kernel.
- Generated a comprehensive diff file after committing the changes to the repository.

2. Work done by Prasun Saxena:

- Tested the machine to check its capability for the VMX virtualization and feature recognition.
- Tested and verified the proper working of the functionality of code by comparing it with the sample output given to us.
- Simulating the answers for the questions in the report.

Question 2: Describe in detail the steps you used to complete the assignment.

Answer:

This assignment focuses on the creation of a Linux kernel module to query various MSRs to determine the virtualization features present in CPU. The module will then report the features it discovers as system message log.

Prerequisites:

- A machine capable of running Linux
- Features for VMX virtualization exposed

Step followed to complete the assignment-

1. Install git

- sudo apt-get install git

2. Download the Linux kernel source code:

- git clone <https://github.com/torvalds/linux.git>

3. Build the Linux Kernel using the source code in previous step:

- sudo apt-get install libncurses-dev
- sudo apt-get install libssl-dev
- make menuconfig

- make
- make modules
- make modules_install
- make install

4. Code a .c (C language) file that will be used to find the capabilities of the MSRs:

-gedit cmpe283-1.c

File content/code displayed in the appendix of this file.

5. Creating new kernel module for MSRs:

-gedit Makefile

File content-

obj-m += cmpe283-1.o

all:

make -C /lib/modules/\$(shell uname -r)/build M=\$(PWD) modules

clean:

make -C /lib/modules/\$(shell uname -r)/build M=\$(PWD) clean

-make all

6. Inserting/Loading the specific kernel module into the kernel:

- insmod ./cmpe283-1.ko

7. Verify the message buffer/output from the kernel in the system message log:

- dmesg

8. Commit the changes done to the local repository:

- git add /home/prasun/Desktop/linux/anshul/cmpe283-1.c
/home/prasun/Desktop/linux/anshul/Makefile

- git config --global user.email "anshulpatni19@gmail.com"

- git config --global user.name "AnshulPatni"

- git commit -m "CMPE 283 Assignment 1 - Initial Commit"

9. Creation of a diff file:

- git diff HEAD~1 > cmpe283-1.diff

10. Check the git log to find the details of the most recent commit:

- git log

```
commit 487e2a060765ea198eb9c66744c5799f13120876
Author: AnshulPatni <anshulpatni19@gmail.com>
Date: Sun Mar 18 23:11:43 2018 -0700
```

CMPE 283 Assignment 1 - Initial commit

```
commit 6f7da290413ba713f0cdd9ff1a2a9bb129ef4f6c
Author: Linus Torvalds <torvalds@linux-foundation.org>
Date: Sun Jul 2 16:07:02 2017 -0700
```

Linux 4.12

```
commit 401e000ab90d7b81d8ea0735e3ff909548754876
Author: Sylvain 'ythier' Hitier <sylvain.hitier@gmail.com>
Date: Sun Jul 2 15:21:56 2017 +0200
```

moduleparam: fix doc: hwparam_irq configures an IRQ

Signed-off-by: Sylvain 'ythier' Hitier <sylvain.hitier@gmail.com>
Signed-off-by: Linus Torvalds <torvalds@linux-foundation.org>

Appendix

1. Code

```
/*
 * cmpe283-1.c - Kernel module for CMPE283 assignment 1
 */
#include <linux/module.h> /* Needed by all modules */
#include <linux/kernel.h> /* Needed for KERN_INFO */
#include <asm/msr.h>

#define MAX_MSG 80

/*
 * Model specific registers (MSRs) by the module.
 * See SDM volume 4, section 2.1
 */
#define IA32_VMX_PINBASED_CTLS 0x481

#define IA32_VMX_BASIC_CTLS 0x480

#define IA32_VMX_PROCBASED_CTLS 0x482
```

```

#define IA32_VMX_PROCBASED_CTLSS2 0x48B

#define IA32_VMX_EXIT_CTLSS 0x483

#define IA32_VMX_ENTRY_CTLSS 0x484

#define IA32_VMX_TRUE_PINBASED_CTLSS 0x48D

#define IA32_VMX_TRUE_PROCBASED_CTLSS 0x48E

#define IA32_VMX_TRUE_EXIT_CTLSS 0x48F

#define IA32_VMX_TRUE_ENTRY_CTLSS 0x490

```

```

/*
 * struct capability_info
 *
 * Represents a single capability (bit number and description).
 * Used by report_capability to output VMX capabilities.
 */

```

```

struct capability_info {
    uint8_t bit;
    const char *name;
};

```

```

/*
 * Pinbased capabilities
 * See SDM volume 3, section 24.6.1
 */
struct capability_info pinbased[5] =
{
    { 0, "External Interrupt Exiting" },
    { 3, "NMI Exiting" },
    { 5, "Virtual NMIs" },
    { 6, "Activate VMX Preemption Timer" },
    { 7, "Process Posted Interrupts" }
};

```

```

/*
 * Procbased Controls
 */
struct capability_info procbased[21] =
{
    { 2, "Interrupt-window" },
    { 3, "Use TSC offsetting" },
    { 7, "HLT exiting" },
    { 9, "INVLPG exiting" },

```

```

    {10, "MWAIT exiting" },
    {11, "RDPMC exiting" },
    {12, "RDTSC exiting" },
    {15, "CR3-load exiting" },
    {16, "CR3-store exiting" },
    {19, "CR8-load exiting" },
    {20, "CR8-store exiting" },
    {21, "Use TPR shadow" },
    {22, "NMI-window exiting" },
    {23, "MOV-DR exiting" },
    {24, "Unconditional I/O" },
    {25, "Use I/O bitmaps" },
    {27, "Monitor trap flag" },
    {28, "Use MSR Bitmaps" },
    {29, "MONITOR exiting" },
    {30, "PAUSE exiting" },
    {31, "Activate secondary controls" }
};

/*
 * Procbased_ctls2 Controls
 */
struct capability_info procbased_ctls2[23] =
{
    { 0, "Virtualize APIC accesses" },
    { 1, "Enable EPT" },
    { 2, "Descriptor-table exiting" },
    { 3, "Enable RDTSCP" },
    { 4, "Virtualize x2APIC mode" },
    { 5, "Enable VPID" },
    { 6, "WBINVD exiting" },
    { 7, "Unrestricted guest" },
    { 8, "APIC-register virtualization" },
    { 9, "Virtual-interrupt delivery" },
    { 10, "PAUSE-loop exiting" },
    { 11, "RDRAND exiting" },
    { 12, "Enable INVPCID" },
    { 13, "Enable VM functions" },
    { 14, "VMCS shadowing" },
    { 15, "Enable ENCLS exiting" },
    { 16, "RDSEED exiting" },
    { 17, "Enable PML" },
    { 18, "EPT-violation" },
    { 19, "Conceal VMX nonroot operation from Intel PT" },
    { 20, "Enable XSAVES/XRSTORS" },
    { 22, "Mode-based execute control for EPT" },
    { 25, "Use TSC scaling" }
};

```

```

/*
 * Entry Controls
 */
struct capability_info entry[9] =
{
    { 2, "Load debug controls" },
    { 9, "IA-32e mode guest" },
    { 10, "Entry to SMM" },
    { 11, "Deactivate dual-monitor treatment" },
    { 13, "Load IA32_PERF_GLOBAL_CTRL" },
    { 14, "Load IA32_PAT" },
    { 15, "Load IA32_EFER" },
    { 16, "Load IA32_BNDCFGS" },
    { 17, "Conceal VM entries from Intel PT" }
};

/*
 * Exit Controls
 */
struct capability_info exit1[11] =
{
    { 2, "Save debug controls" },
    { 9, "Host addressspace size" },
    { 12, "Load IA32_PERF_GLOBAL_CTRL" },
    { 15, "Acknowledge interrupt on exit" },
    { 18, "Save IA32_PAT" },
    { 19, "Load IA32_PAT" },
    { 20, "Save IA32_EEFR" },
    { 21, "Load IA32_EFER" },
    { 22, "Save VMXpreemption timer value" },
    { 23, "Clear IA32_BNDCFGS" },
    { 24, "Conceal VM exits from Intel PT" },
};

/*
 * report_capability
 *
 * Reports capabilities present in 'cap' using the corresponding MSR values
 * provided in 'lo' and 'hi'.
 *
 * Parameters:
 * cap: capability_info structure for this feature
 * len: number of entries in 'cap'
 * lo: low 32 bits of capability MSR value describing this feature
 * hi: high 32 bits of capability MSR value describing this feature
 */

```

```

void
report_capability(struct capability_info *cap, uint8_t len, uint32_t lo,
uint32_t hi)
{
    uint8_t i;
    struct capability_info *c;
    char msg[MAX_MSG];

    memset(msg, 0, sizeof(msg));

    for (i = 0; i < len; i++) {
        c = &cap[i];
        snprintf(msg, 79, " %s: Can set=%s, Can clear=%s\n",
            c->name,
            (hi & (1 << c->bit)) ? "Yes" : "No",
            !(lo & (1 << c->bit)) ? "Yes" : "No");
        printk(msg);
    }
}

```

```

/*
 * detect_vmx_features
 *
 * Detects and prints VMX capabilities of this host's CPU.
 */

```

```

void
detect_vmx_features(void)
{
    uint32_t lo, hi;

    rdmsr(IA32_VMX_BASIC_CTLs, lo, hi);
    pr_info("Basic Controls MSR: 0x%llx\n",
        (uint64_t)(lo | (uint64_t)hi << 32));
    //report_capability(pinbased, 5, lo, hi);

    pr_info("\n");

    if(!((hi>>23)&& 1))
    {

        pr_info("True Controls capability is not set");

        pr_info("\n");

        /* Pinbased controls */
        rdmsr(IA32_VMX_PINBASED_CTLs, lo, hi);
        pr_info("Pinbased Controls MSR: 0x%llx\n",
            (uint64_t)(lo | (uint64_t)hi << 32));
        report_capability(pinbased, 5, lo, hi);
    }
}

```

```

pr_info("\n");

/* Procbased controls */
rdmsr(IA32_VMX_PROCBASED_CTL, lo, hi);
pr_info("Procbased Controls MSR: 0x%llx\n",
        (uint64_t)(lo | (uint64_t)hi << 32));
report_capability(procbased, 21, lo, hi);

pr_info("\n");

if((hi>>31)&& 1)
{
    /* Procbased_ctls2 controls */
    rdmsr(IA32_VMX_PROCBASED_CTL2, lo, hi);
    pr_info("Procbased_ctls2 Controls MSR: 0x%llx\n",
            (uint64_t)(lo | (uint64_t)hi << 32));
    report_capability(procbased_ctls2, 23, lo, hi);
}
else
{
    pr_info("No capability set for “Activate Secondary Controls” control in the
primary procbased controls");
}

pr_info("\n");

/* Exit controls */
rdmsr(IA32_VMX_EXIT_CTL, lo, hi);
pr_info("Exit Controls MSR: 0x%llx\n",
        (uint64_t)(lo | (uint64_t)hi << 32));
report_capability(exit1, 11, lo, hi);

pr_info("\n");

/* Entry controls */
rdmsr(IA32_VMX_ENTRY_CTL, lo, hi);
pr_info("Entry Controls MSR: 0x%llx\n",
        (uint64_t)(lo | (uint64_t)hi << 32));
report_capability(entry, 9, lo, hi);
}
else
{

pr_info("True Controls capability is set");

pr_info("\n");

```



```

        /* True Pinbased controls */
        rdmsr(IA32_VMX_TRUE_PINBASED_CTL, lo, hi);
        pr_info("True Pinbased Controls MSR: 0x%llx\n",
                (uint64_t)(lo | (uint64_t)hi << 32));
        report_capability(pinbased, 5, lo, hi);

        pr_info("\n");

        /* True Procbased controls */
        rdmsr(IA32_VMX_TRUE_PROCBASED_CTL, lo, hi);
        pr_info("True Procbased Controls MSR: 0x%llx\n",
                (uint64_t)(lo | (uint64_t)hi << 32));
        report_capability(procbased, 21, lo, hi);

        pr_info("\n");

        /* True Exit controls */
        rdmsr(IA32_VMX_TRUE_EXIT_CTL, lo, hi);
        pr_info("True Exit Controls MSR: 0x%llx\n",
                (uint64_t)(lo | (uint64_t)hi << 32));
        report_capability(exit1, 11, lo, hi);

        pr_info("\n");

        /* True Entry controls */
        rdmsr(IA32_VMX_TRUE_ENTRY_CTL, lo, hi);
        pr_info("True Entry Controls MSR: 0x%llx\n",
                (uint64_t)(lo | (uint64_t)hi << 32));
        report_capability(entry, 9, lo, hi);
    }

}

/*
 * init_module
 *
 * Module entry point
 *
 * Return Values:
 * Always 0
 */
int
init_module(void)
{
    printk(KERN_INFO "CMPE 283 Assignment 1 Module Start\n");

    detect_vmx_features();

```

```

    /*
    * A non 0 return means init_module failed; module can't be loaded.
    */
    cleanup_module();
    return 0;
}

/*
* cleanup_module
*
* Function called on module unload
*/
void
cleanup_module(void)
{
    printk(KERN_INFO "CMPE 283 Assignment 1 Module Exits\n");
}

```

Output:

```

[ 9517.403597] CMPE 283 Assignment 1 Module Start
[ 9517.403602] Basic Controls MSR: 0xda040000000012

[ 9517.403605] True Controls capability is set

[ 9517.403610] True Pinbased Controls MSR: 0x7f00000016
[ 9517.403613] External Interrupt Exiting: Can set=Yes, Can clear=Yes
[ 9517.403617] NMI Exiting: Can set=Yes, Can clear=Yes
[ 9517.403620] Virtual NMIs: Can set=Yes, Can clear=Yes
[ 9517.403623] Activate VMX Preemption Timer: Can set=Yes, Can clear=Yes
[ 9517.403626] Process Posted Interrupts: Can set=No, Can clear=Yes

[ 9517.403630] True Procbased Controls MSR: 0xffff9fffe04006172
[ 9517.403634] Interrupt-window: Can set=Yes, Can clear=Yes
[ 9517.403637] Use TSC offsetting: Can set=Yes, Can clear=Yes
[ 9517.403640] HLT exiting: Can set=Yes, Can clear=Yes
[ 9517.403643] INVLPG exiting: Can set=Yes, Can clear=Yes
[ 9517.403645] MWAIT exiting: Can set=Yes, Can clear=Yes
[ 9517.403648] RDPMC exiting: Can set=Yes, Can clear=Yes
[ 9517.403651] RDTSC exiting: Can set=Yes, Can clear=Yes
[ 9517.403654] CR3-load exiting: Can set=Yes, Can clear=Yes
[ 9517.403657] CR3-store exiting: Can set=Yes, Can clear=Yes
[ 9517.403661] CR8-load exiting: Can set=Yes, Can clear=Yes
[ 9517.403664] CR8-store exiting: Can set=Yes, Can clear=Yes
[ 9517.403667] Use TPR shadow: Can set=Yes, Can clear=Yes
[ 9517.403670] NMI-window exiting: Can set=Yes, Can clear=Yes
[ 9517.403673] MOV-DR exiting: Can set=Yes, Can clear=Yes
[ 9517.403676] Unconditional I/O: Can set=Yes, Can clear=Yes

```

[9517.403680] Use I/O bitmaps: Can set=Yes, Can clear=Yes
[9517.403683] Monitor trap flag: Can set=Yes, Can clear=Yes
[9517.403686] Use MSR Bitmaps: Can set=Yes, Can clear=Yes
[9517.403689] MONITOR exiting: Can set=Yes, Can clear=Yes
[9517.403692] PAUSE exiting: Can set=Yes, Can clear=Yes
[9517.403696] Activate secondary controls: Can set=Yes, Can clear=Yes

[9517.403700] True Exit Controls MSR: 0x7ffff00036dfb
[9517.403703] Save debug controls: Can set=Yes, Can clear=Yes
[9517.403707] Host addressspace size: Can set=Yes, Can clear=Yes
[9517.403710] Load IA32_PERF_GLOBAL_CTRL: Can set=Yes, Can clear=Yes
[9517.403714] Acknowledge interrupt on exit: Can set=Yes, Can clear=Yes
[9517.403717] Save IA32_PAT: Can set=Yes, Can clear=Yes
[9517.403720] Load IA32_PAT: Can set=Yes, Can clear=Yes
[9517.403723] Save IA32_EEFR: Can set=Yes, Can clear=Yes
[9517.403750] Load IA32_EFER: Can set=Yes, Can clear=Yes
[9517.403765] Save VMXpreemption timer value: Can set=Yes, Can clear=Yes
[9517.403774] Clear IA32_BNDCFGS: Can set=No, Can clear=Yes
[9517.403787] Conceal VM exits from Intel PT: Can set=No, Can clear=Yes

[9517.403817] True Entry Controls MSR: 0xffff000011fb
[9517.403830] Load debug controls: Can set=Yes, Can clear=Yes
[9517.403842] IA-32e mode guest: Can set=Yes, Can clear=Yes
[9517.403855] Entry to SMM: Can set=Yes, Can clear=Yes
[9517.403869] Deactivate dual-monitor treatment: Can set=Yes, Can clear=Yes
[9517.403881] Load IA32_PERF_GLOBAL_CTRL: Can set=Yes, Can clear=Yes
[9517.403898] Load IA32_PAT: Can set=Yes, Can clear=Yes
[9517.403907] Load IA32_EFER: Can set=Yes, Can clear=Yes
[9517.403919] Load IA32_BNDCFGS: Can set=No, Can clear=Yes
[9517.403931] Conceal VM entries from Intel PT: Can set=No, Can clear=Yes
[9517.403944] CMPE 283 Assignment 1 Module Exits